

Zero-Abstraction Systems: Architectural Scaling of a Lock-Free Distributed Single-Level Storage Operating System over Native NVMe Queues

Dave F. Cook
Independent Researcher
dave@gridworkz.com

ABSTRACT

Modern computer systems remain bound by legacy abstractions that isolate volatile runtime state from non-volatile persistent storage. The traditional Virtual File System (VFS) layer, while generic, introduces critical inefficiencies including redundant memory double-caching, expensive file-descriptor serialization loops, and deep software kernel stacks that struggle to keep pace with microsecond-level NVMe PCIe Gen5 solid-state media.

This paper introduces *AeroSLS*, a custom-built, multi-core 64-bit operating system implementing an un-abstracted Single-Level Storage (SLS) architecture. AeroSLS eliminates the file concept entirely, representing all persistent data blocks as a unified, globally addressable 64-bit object namespace mapped directly by the processor’s Memory Management Unit (MMU).

To ensure strict safety and multi-core scaling under parallel load, AeroSLS implements a lock-free, concurrent FNV-1a object directory, a priority-driven hardware NVMe I/O queue traffic broker, and parallelized AVX-512 cryptographic page-sealing engines. Furthermore, we mitigate cluster split-brain anomalies and hardware execution jitter using an in-kernel Raft-lite consensus state machine paired with a lazy floating-point context-switching framework.

Evaluations conducted in a four-core QEMU environment show that AeroSLS’s object-lookup step averages 17.2 cycles per resolution event (a 96.1% reduction versus a simulated VFS path), and that the lazy CR0.TS context-switch policy reduces per-switch cycles by 54.2× relative to forced full-state saves (Section 6).

KEYWORDS

Single-Level Storage, Persistent Memory, Operating Systems, NVMe, AVX-512, Lock-Free, Distributed Systems

ACM Reference Format:

Dave F. Cook. 2026. Zero-Abstraction Systems: Architectural Scaling of a Lock-Free Distributed Single-Level Storage Operating System over Native NVMe Queues. In *SysArch ’26: Symposium*

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

SysArch ’26, 2026, Virtual

© 2026 Copyright held by the owner/author(s).

ACM ISBN 978-1-4503-XXXX-X/26/01

<https://doi.org/XXXXXXX.XXXXXXX>




Figure 1: Architectural comparison: legacy VFS double-caching stack vs. the AeroSLS direct hardware memory-mapped persistence layout.

on *Operating Systems Architecture*, 2026. ACM, New York, NY, USA, 7 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 INTRODUCTION

For over half a century, operating system design has separated volatile runtime state from non-volatile persistent storage. This architectural dualism forces applications to operate across two disparate mental models: a granular byte-addressable space managed by high-speed memory architectures, and a block-oriented file abstraction managed via systemic block I/O boundary layers.

To access persistent blocks, application pipelines must navigate the complete Virtual File System (VFS) layers, endure data parsing marshaling overheads, and allocate localized user-space buffers. While historically necessary to mask slow mechanical disk rotations, this model imposes a high software abstraction cost when applied to ultra-low-latency modern flash cells and byte-addressable persistent architectures.

With native NVMe solid-state storage devices operating directly across high-speed PCIe lanes, hardware processing latencies have fallen to single-digit microseconds. Under these conditions, the dominant system bottleneck shifts from physical hardware transport delays to the software operating system itself. Deep, synchronous kernel storage stacks, global filesystem metadata lock contentions, and internal memory cache duplication degrade total pipeline performance [11].

To bypass these abstraction layers, we introduce *AeroSLS*, a clean-slate, distributed, multi-core 64-bit operating system designed to fully converge primary volatile execution

memory and permanent block storage. AeroSLS takes a zero-abstraction approach: it removes files, directories, mounts, and block allocation maps from the kernel. Instead, all data blocks throughout the machine sit directly within a unified, globally available 64-bit virtual memory address space.

When a program alters a structure pointer, the CPU's underlying MMU directly tracks the transaction. If a targeted memory segment resides natively on persistent flash media blocks instead of volatile RAM, the processor fires a hardware Page Fault exception. AeroSLS traps this signal, evaluates the underlying target space via an internal lock-free concurrent hash index matrix, and schedules an asynchronous direct DMA payload write straight onto native NVMe Submission and Completion Ring queues with zero intermediate buffers or user-space serialization loops.

This paper makes the following primary contributions:

- We design an un-abstracted SLS operating system architecture that unifies memory and persistent media at the processor page-table boundary, removing file, directory, and mount abstractions from the kernel (Sections 2–3).
- We implement a multi-core I/O pipeline combining lock-free FNV-1a directory indices, priority-driven per-core NVMe queue pairs, and MSI-X interrupt routing via Local APIC (Sections 2–3).
- We introduce an AVX-512 vectorized ChaCha20 page-sealing engine paired with a lazy CR0.TS context switcher that reduces per-switch jitter without sacrificing data-at-rest confidentiality (Section 4).
- We measure two principal overheads in a four-core QEMU environment: object-lookup resolution (17.2 mean cycles vs. 446.5 for a simulated VFS path, a 96.1% reduction) and scheduler context-switch cost (45.8 mean cycles under the lazy policy vs. 2,485.0 under forced full-state saves, a 54.2× improvement) (Section 6).

The rest of this paper is structured as follows. Section 2 details our unified address mapping matrix and concurrent index architectures. Section 3 describes the hardware integration and multi-core scalability mechanisms. Section 4 presents the memory protection and data privacy subsystems. Section 5 covers crash consistency and distributed fault tolerance. Section 6 presents our experimental evaluation. Section 7 surveys related work, and Section 8 concludes.

2 SYSTEM ARCHITECTURE DESIGN

The architectural layout of AeroSLS departs entirely from the traditional layered virtual directory design found in monolithic kernels. By implementing a flat address tracking strategy, the operating system converts storage blocks into standard architecture page structures. This removes the storage translation and caching overhead imposed by a Virtual File System (VFS).

2.1 Unified 64-bit Address Topology

AeroSLS splits the x86_64 48-bit canonical linear space into two distinct hardware privilege domains. The lower memory region is allocated strictly to transient scheduler threads, kernel text blocks, and processor descriptors. The upper address range, anchored securely at the linear boundary address pointer base 0x0000700000000000, forms the global Single-Level Storage unified address space.

Every individual file, table, or workspace segment allocated by a user shell instance is mapped directly to a distinct 64-bit virtual memory coordinate window within this domain via the processor's Page Map Level 4 (PML4) allocation frames. Because data spaces map to permanent unique identifiers, this structural memory layout remains uniform across both volatile primary memory (RAM) and non-volatile secondary storage (NVMe flash sectors).

2.2 Concurrent Lock-Free Object Indexing

To prevent multi-core execution stalls when thousands of background threads issue concurrent lookups via system call gates, AeroSLS implements a concurrent chained lookup matrix index. This design replaces global kernel mutex chains with atomic compare-and-swap (CAS) primitives.

Algorithm 1 Lock-Free Head Node Entry Insertion

Require: *Object_ID*, *Virtual_Address*, *Bucket_Index*

```

1: Node ← AllocateKernelMemory(sizeof(SLSObjectNode))

2: Node.ID ← Object_ID
3: Node.VAddr ← Virtual_Address
4: loop
5:   Current_Head ← GlobalHashMatrix[Bucket_Index]
6:   Node.Next ← Current_Head
7:   if CompareAndSwap(&GlobalHashMatrix[Bucket_Index], Current_Head, Node)
   then
8:     break {Node committed without locking}
9:   end if
10: end loop

```

As formalized in Algorithm 1, if a race condition occurs between Core 0 and Core 1 attempting an insertion into the same bucket tracking row, the **CompareAndSwap** function handles the serialization atomically at the hardware memory controller layer, avoiding global context yields or thread stall states.

2.3 Linear Extent Translation Matrix

To prevent physical block device fragmentation from breaking contiguous virtual memory spaces, AeroSLS decouples linear virtual allocations from underlying storage sectors using a flat *Extent Mapping Table*. When an object is created or dynamically scaled, the system records it as a collection of non-contiguous physical disk chunks called extents. An individual extent is mathematically modeled as:

$$\mathcal{E} = \langle V_{page}, L_{lba}, S_{count} \rangle \quad (1)$$

where V_{page} represents the starting virtual page offset index within the object container, L_{lba} represents the destination physical Logical Block Address on the NVMe device, and S_{count} indicates the total contiguous sector block count length.

When a thread triggers a page fault exception on a non-resident linear virtual memory coordinate ($\mathcal{V}_{fault}$), the kernel page-fault engine resolves the correct sector mapping using the target offset evaluation function:

$$\mathcal{L}_{target} = L_{lba} + \left(\left\lfloor \frac{\mathcal{V}_{fault} - \mathcal{V}_{base}}{4096} \right\rfloor \times 8 - V_{page} \times 8 \right) \quad (2)$$

where $\mathcal{V}_{base}$ is the universal root allocation base pointer of the target memory object. This calculation resolves the fragmented block offset in constant time, bypassing traditional multi-tier filesystem node traversal.

2.4 Predictive Spatial Pre-Fetching Architecture

To minimize network and media access latencies, AeroSLS utilizes a predictive spatial pre-fetching engine. When an application thread touches page $\mathcal{P}_n$, causing a hardware page fault interrupt, the engine analyzes the traversal pattern. If sequential execution is detected, the kernel generates asynchronous pre-fetch operations for pages $\mathcal{P}_{n+1}$ and $\mathcal{P}_{n+2}$.

These predictive reads are dispatched via low-priority non-blocking commands across parallel PCIe or Ethernet channels. The incoming data blocks are packed into allocated physical memory frames and updated within the processor page tables before the user thread advances its code pointer, converting predictable storage accesses into hardware RAM cache hits before the user thread advances.

3 HARDWARE INTEGRATION AND MULTI-CORE SCALABILITY

3.1 NVMe Submission and Completion Queue Architecture

AeroSLS communicates with persistent storage exclusively through native NVMe hardware queues, bypassing all block layer and SCSI translation intermediaries [8]. At boot, the kernel's PCI bus scanner enumerates configuration space registers to locate the NVMe controller's Memory-Mapped I/O (MMIO) Base Address Register (BAR0).

Once discovered, the driver initializes an *Admin Queue Pair* (Submission Queue / Completion Queue depth 64) to issue controller configuration commands, followed by one I/O Queue Pair per logical CPU core. This per-core queue topology removes cross-core lock contention on the submission path: a page-fault handler running on Core k deposits a 64-byte Submission Queue Entry (SQE) directly into its dedicated queue, rings the hardware doorbell register, and returns immediately to userspace.

The NVMe controller signals completion by writing a 16-byte Completion Queue Entry (CQE) and asserting the configured MSI-X interrupt vector, waking the waiting kernel thread without polling. Priority differentiation between

latency-sensitive page-fault I/Os and background flush operations is enforced by the in-kernel *I/O Priority Traffic Broker*, which selectively populates high- and low-priority queue slots according to real-time queue depth telemetry.

3.2 Symmetric Multiprocessing Bootstrap

AeroSLS wakes Application Processors (APs) using a 16-bit real-mode trampoline payload injected at physical address `0x08000` [6]. The Bootstrap Processor (BSP) copies the binary blob into the low-memory page, writes the shared PML4 base pointer into the metadata handshake vector at `0x07000`, and dispatches an INIT/SIPI interrupt sequence to each AP's Local APIC.

Each woken AP executes the trampoline, enables Physical Address Extension (PAE), clones the BSP's PML4 page table, and transitions atomically to 64-bit Long Mode. A per-AP stack pointer is pre-allocated from the physical frame pool and written into the handshake vector at `0x07030`. Once the AP acquires its unique stack, it sets the `ap.bootstrap_lock` atomic flag to signal the BSP, which advances to the next processor, serializing the bring-up sequence safely.

3.3 Distributed Interrupt Routing via APIC

PCIe device completion interrupts are distributed across application cores using the Local APIC and I/O APIC redirection matrices. The NVMe controller's MSI-X capability table is programmed so that queue k completion vectors route to core k 's Local APIC. This ensures that the interrupt handler executing on core k can directly update the page table and wake the blocked thread without IPI cross-calls, eliminating inter-core cache-line synchronization overhead on the critical path.

4 MEMORY PROTECTION AND DATA PRIVACY AT REST

4.1 Hardware-Enforced Capability Access Control

AeroSLS enforces object-level access permissions at the Page Table Entry (PTE) level rather than through software permission string evaluation paths [6]. The *User/Supervisor* flag (bit 2) in each PTE determines whether a given object page is accessible from Ring 3. When an object is created, the kernel assigns a permission bitmap and configures the corresponding PTE flags accordingly. All subsequent permission checks are executed by the hardware MMU at no additional instruction cost; security enforcement is structurally embedded in the address translation path.

4.2 Parallel AVX-512 Cryptographic Page Sealing

To ensure data-at-rest confidentiality across persistent NVMe sectors, AeroSLS dedicates Cores 2 and 3 as cryptoprocessors running a vectorized ChaCha20 stream cipher [7]. When the dirty-page flush daemon schedules a frame for write-back, the

Core Allocation Topology




Figure 2: Core allocation topology: Cores 0–1 handle scheduler threads and allocations; Cores 2–3 function as isolated vectorized ChaCha20 cryptoprocessors.

frame is dispatched to the crypto pipeline rather than the NVMe queue directly. The cipher operates on 64-byte aligned AVX-512 register blocks (`zmm0–zmm31`), sealing the 4 KiB page in a single vectorized pass before the DMA transfer is submitted.

Because NVMe page frames are sourced from the physical frame pool, they are already 4 KiB-aligned, satisfying the 64-byte boundary required by `vmovdqa64` without additional padding. Violations are caught by an alignment assertion in the assembly prologue, which traps to the fault handler before any out-of-bounds vector operation can trigger an unrecoverable General Protection Fault (Exception 13).

4.3 Lazy Floating-Point Context Switching

When a standard integer thread is context-switched, naively saving the full 2,688-byte AVX-512 extended state on every rotation imposes a prohibitive penalty. AeroSLS implements a *Lazy Floating-Point Save* optimization using the hardware `CR0.TS` trap mechanism.

State size model. The total physical byte payload required by a strict context switch $\mathcal{S}_{total}$ is:

$$\mathcal{S}_{total} = \mathcal{S}_{legacy} + \mathcal{S}_{avx2} + \mathcal{S}_{avx512ext} + \mathcal{S}_{alignment} \quad (3)$$

where $\mathcal{S}_{legacy} = 512$ B (x87/SSE `fxsave` footprint), $\mathcal{S}_{avx2} = 256$ B (upper 128-bit halves of `ymm0–ymm15`), and $\mathcal{S}_{avx512ext} = 512 + 1024 + 64 = 1600$ B (upper ZMM extensions, `zmm16–zmm31`, and `opmask` registers `k0–k7`). Substituting the fixed hardware constants:

$$\mathcal{S}_{total} = 512 + 256 + 1600 + 320 = \mathbf{2688 \text{ B}} \quad (4)$$

Latency model. Let $\mathcal{C}_{gpr}$ be the cost of saving general-purpose integer registers, $\mathcal{C}_{xsave}$ the multi-cycle cost of moving the 2,688-byte vector block across internal memory buses, and $\mathcal{U}_{vector} \in \{0, 1\}$ a binary variable indicating whether the active thread executes vector instructions.

Under a **strict switching policy**:

$$\mathcal{L}_{strict} = \mathcal{C}_{gpr} + \mathcal{C}_{xsave} + \mathcal{C}_{xrstor} \quad (5)$$

Under the **lazy CR0.TS optimization**:

$$\mathcal{L}_{lazy} = \mathcal{C}_{gpr} + \mathcal{U}_{vector} \cdot (\mathcal{C}_{trap} + \mathcal{C}_{xsave} + \mathcal{C}_{xrstor}) \quad (6)$$

where $\mathcal{C}_{trap}$ is the single-digit-cycle cost of the *Device Not Available* Exception 7 handler.

For non-vector threads ($\mathcal{U}_{vector} = 0$):

$$\lim_{\mathcal{U}_{vector} \rightarrow 0} \mathcal{L}_{lazy} = \mathcal{C}_{gpr} \quad (7)$$

This limit formalizes the evaluation finding in Section 6: for scalar threads, AeroSLS removes the 2,688-byte vector-state save from the context-switch path.

5 CRASH CONSISTENCY AND DISTRIBUTED FAULT TOLERANCE

5.1 Write-Ahead Journaling and Object Directory

AeroSLS maintains crash consistency through a two-tier on-disk layout. Sector 1024 anchors the *Global Object Directory* (GOD), identified by the eight-byte magic token `SLSROOTD`. Sector 2048 hosts the *Write-Ahead Journal* (WAJ) ledger, a micro-commit log where pending mutations are recorded before any corresponding page frame is modified. On cold boot, the kernel validates the GOD magic and replays any uncommitted WAJ entries, restoring the object namespace to the last consistent state prior to any power loss.

5.2 In-Kernel Raft-Lite Consensus Engine

When AeroSLS scales across multiple network nodes, a network partition can produce a *split-brain* anomaly where two nodes concurrently believe they hold exclusive write authority over a shared persistent object. To prevent this, Core 3 on each node executes a Raft-lite [9] state machine with three roles: LEADER, FOLLOWER, and CANDIDATE.

The leader broadcasts heartbeat packets over the e1000 network interface every 10 ms. If a follower receives no heartbeat within a 150-tick timeout window (1.5 s), it transitions to CANDIDATE, increments its election term, and immediately invokes `update_page_table_permissions_globally(1)` to strip all locally mapped objects of write authorization. This prevents local writes during the election window, reducing the risk of divergence. Once a candidate accumulates a strict majority of VOTEREPLY acknowledgments, it is elevated to LEADER and write permissions are restored.

5.3 Distributed SLS Page Protocol (DSPP)

AeroSLS replicates dirty memory pages across cluster nodes using the *Distributed SLS Page Protocol* (DSPP), a custom encapsulation layer transmitted over standard Ethernet via the Intel e1000 adapter. Each DSPP frame carries a 48-byte header containing a `SLSNETMA` magic marker, the 64-bit `system_object_id`, the faulting `virtual_address`, an atomic

`transaction_id`, an opcode (`DSPP_PAGE_WRITE_REQ / ACK / PAGE_READ_REQ / ACK`), and a 4 KiB payload frame.

When an application invokes `sys_sls_fence`, the kernel pins the local RAM frame, constructs a full DSPP packet, posts it to the e1000 Transmit Descriptor Ring, and blocks the calling thread on a *network transaction token*. On the remote node, the MSI-X interrupt fires on Core 3, which deserializes the packet, writes the 4 KiB payload into the corresponding local frame, and returns a `DSPP_PAGE_WRITE_ACK`. The originating node matches the transaction ID, unblocks the application thread, and only then advances the WAJ commit pointer.

6 EVALUATION AND EXPERIMENTAL RESULTS

6.1 Experimental Setup

All benchmarks were conducted inside a QEMU 7.2 emulated environment configured with four x86_64 virtual cores, 4 GiB RAM, a virtio-backed NVMe device, and an Intel e1000 virtual network adapter. Cycle measurements used the processor’s `rdtsc` instruction, with an initial warm-up phase of ten iterations discarded to eliminate JIT and cache cold-start outliers. Each reported result is the median of 100 back-to-back trials. Two metric families were evaluated: the *memory resolution abstraction tax* and the *core scheduler context-switch jitter*.

6.2 Memory Resolution Abstraction Tax




Figure 3: Dual-panel evaluation results. (a) Memory resolution abstraction tax: legacy VFS stack versus AeroSLS direct MMU index. (b) Core scheduler jitter mitigation: forced strict XSAVE/XRSTOR versus lazy CR0.TS context switching. Panel (b) uses a logarithmic y -axis.

The legacy VFS simulation traverses a synthetic file-path string, performs a file-descriptor table boundary check, and copies a 512-byte intermediate kernel buffer — reproducing the dominant instruction-count categories found in Linux VFS `open/read` call paths. The AeroSLS path computes an FNV-1a bucket index directly from the 64-bit object ID and returns the corresponding virtual address window. Both paths are measured end-to-end from `rdtsc` entry to `rdtsc` exit with no I/O involved.

As shown in Figure 3(a) and Table 1, the legacy VFS model incurs a mean cost of **446.5 clock cycles** per resolution event. The AeroSLS direct MMU path reduces this to a mean of **17.2 cycles** (a **96.1%** reduction in execution overhead). The AeroSLS distribution also exhibits dramatically tighter variance (min 14, max 22 cycles) versus the legacy stack (min 412, max 498 cycles), consistent with lock-free hash addressing reducing the unpredictable synchronization delays common in monolithic kernels.

6.3 Core Scheduler Context-Switch Jitter

When a cryptographic worker on Core 2 is preempted and an integer shell thread is scheduled in, a strict policy forces a full `xsave/xrstor` cycle moving the 2,688-byte AVX-512 state buffer. The lazy policy instead sets `CR0.TS` and defers the vector state save to the first vector instruction executed by the new thread (Device Not Available Exception 7), if any.

The results in Figure 3(b) show that the strict policy incurs a mean latency of **2,485.0 cycles** per switch. With the lazy optimization active, the mean drops to **45.8 cycles** (a **54.2× improvement**). Because interactive shell threads are wholly scalar, the $U_{vector} = 0$ branch is taken on virtually every switch, consistent with the analytical model in Section 4.3.

Table 1: Per-metric cycle count summary across 100 trial iterations (QEMU x86_64, 4-core, 4 GiB RAM).

Metric	Min	Mean	Max
Legacy VFS Stack Cycles	412	446.5	498
AeroSLS Direct MMU Cycles	14	17.2	22
Lazy CR0.TS Switch Cycles	38	45.8	62
Strict XSAVE/XRSTOR Cycles	2,241	2,485.0	2,731

When mapping the mean AeroSLS MMU cycle score to the baseline 4 GHz virtual clock rate of the emulation platform, each object resolution completes in approximately 4.3 ns, well within the single-digit-microsecond floor set by PCIe Gen5 NVMe access latency. At this scale, the object-lookup step is no longer the dominant bottleneck in the storage pipeline.

7 RELATED WORK

The concept of a Single-Level Storage operating system dates back to pioneer platforms such as the Multics system [3] and the IBM System/38 (later evolving into the IBM i architecture) [5]. These early designs unified the programming abstraction between memory and disk storage. However, they were engineered for an era dominated by slow mechanical hard disks and single-core processors, relying on hardware-dependent segmented memory configurations and deep software abstraction layers that incurred substantial overhead.

Modern non-volatile, low-latency storage media has renewed interest in persistent memory operating system research. Specialized research kernels like Twizzler [1] and BPFS [2] have explored the removal of standard file system

semantics by converting persistent storage targets into cross-addressable memory object spaces. Twizzler implements an object-centric model that uses data-pointer formatting to navigate cross-object links; however, Twizzler operates under user-space runtime environments that still rely on underlying traditional file systems or hypervisor interventions to coordinate raw media writes.

Clean-slate persistent frameworks such as PMFS [4] and NOVA [10] provide highly optimized, byte-addressable file systems specifically customized for non-volatile dual in-line memory modules (NVDIMMs). While these systems achieve significant throughput gains by stripping legacy block layer drivers from the software stack, they fundamentally preserve the classic POSIX file API boundaries (`open`, `read`, `write`, `close`). Consequently, application pipelines must continue to bear the strict serialization and buffer management taxes required to marshal state across application memory boundaries.

Storage development kits such as SPDK [11] eliminate intermediate kernel I/O stacks via user-space NVMe poll-mode drivers, approaching hardware line rates. However, SPDK preserves the file/object boundary at the application level and does not address the MMU-level convergence of memory and persistent storage that AeroSLS achieves.

AeroSLS differs from these prior frameworks by enforcing a zero-abstraction model at the bare-metal kernel layer. It removes file layers entirely and avoids runtime address-patching overhead, relying directly on native x86.64 hardware paging mechanisms, lock-free directories, and prioritized NVMe hardware queues to achieve crash-consistent data persistence at hardware line rates.

8 CONCLUSION AND FUTURE WORK

This paper presented *AeroSLS*, a multi-core 64-bit operating system designed to eliminate the legacy division between volatile primary memory and non-volatile block storage. By converging storage execution pathways directly at the processor page-table boundary, AeroSLS bypasses the virtualization and double-caching overhead imposed by classic Virtual File System (VFS) layers.

Measurements show that lock-free concurrent hash matrices reduce the object-lookup cost by 96.1% relative to a simulated VFS path, and the lazy CR0.TS context-switch policy cuts per-switch cycles by 54.2× for scalar threads. By isolating cryptographic sealing on dedicated cores and deferring vector-state saves via CR0.TS, the kernel maintains data-at-rest confidentiality without introducing jitter to scalar compute threads.

8.1 Heterogeneous Architecture and Protocol Scaling

While the current instantiation of the Distributed SLS Page Protocol (DSPP) assumes homogeneous hardware profiles across the compute cluster, future work will focus on expanding the protocol into a highly heterogeneous network architecture. Scaling an un-abstracted single-level storage

space across mixed-architecture topologies (e.g., co-linking x86.64, ARM64, and RISC-V nodes) introduces significant structural challenges, most notably variable processor page sizes (4 KiB, 16 KiB, and 64 KiB) and conflicting endianness geometries.

To mitigate these barriers without introducing virtualization taxes, we propose the design of a hardware-agnostic *Page-Slicing Translation Layer* within the kernel's network routing framework. This sub-system will dynamically partition larger remote memory blocks into aligned sub-page extents on demand, converting architecture-specific memory attributes natively at the packet boundary to enable consistent cross-platform object sharing.

8.2 Hardware-Accelerated RDMA over Converged Ethernet

To further reduce remote page-fault resolution latencies, we plan to transition network communications from legacy emulated Ethernet adapters to hardware-accelerated Remote Direct Memory Access (RDMA) over Converged Ethernet (RoCEv2). By interfacing directly with RoCEv2-enabled Network Interface Cards (rNICs) via memory-mapped Queue Pairs, the page-fault handler will bypass the local and remote CPU cores entirely during remote reads.




Figure 4: Proposed AeroSLS zero-copy RoCEv2 integration pipeline for sub-microsecond distributed page resolution.

When a remote page fault triggers, the kernel will post an RDMA Read Request straight to the hardware rNIC PCIe lanes. The remote network card will pull the target 4 KiB data frame directly out of its local physical RAM pool using peer-to-peer DMA over the wire. This pipeline is projected to reduce distributed memory synchronization delays to sub-microsecond scales, approaching the latency of unified hardware memory fabrics (Figure 4).

8.3 SLS-Aware Compiler Plugins for Legacy Runtimes

Finally, we aim to eliminate the necessity for explicit system allocation wrappers by developing automated, compiler-level optimization plugins for legacy programming languages (such as C, C++, and Rust). Interacting with the persistent object heap via manual system call bindings shifts the burden of structural safety checking entirely onto the application developer.

To make the single-level object-heap infrastructure entirely seamless, we are engineering custom compiler front-end

and back-end passes using the LLVM infrastructure. These plugins will intercept standard language keywords (e.g., overloading the C++ `new` operator or introducing a custom Rust attribute macro such as `#[persistent]`) during the Intermediate Representation (LLVM-IR) compilation phase. The compiler will automatically insert offset-based `sls_malloc` hooks, manage relative pointer calculations within individual data structures, and handle automatic alignment constraints at compile time. This allows legacy applications to be ported into the AeroSLS environment and achieve crash-consistent persistence without modifying their core application logic.

REFERENCES

- [1] Daniel Bittman, Peter Alvaro, Pankaj Mehra, Darrell D. E. Long, and Ethan L. Miller. 2020. Twizzler: a Data-Centric OS for Non-Volatile Memory. In *2020 USENIX Annual Technical Conference (USENIX ATC '20)*. USENIX Association, 65–80.
- [2] Jeremy Condit, Edmund B. Nightingale, Christopher Frost, Engin Ipek, Benjamin C. Lee, Doug Burger, and Derrick Coetzee. 2009. Better I/O Through Byte-Addressable, Persistent Memory. In *Proceedings of the 22nd ACM SIGOPS Symposium on Operating Systems Principles (SOSP '09)*. ACM, 133–146.
- [3] Fernando J. Corbató and Victor A. Vyssotsky. 1965. Introduction and Overview of the Multics System. In *Proceedings of the Fall Joint Computer Conference (FJCC '65)*. ACM/IEEE, 185–196.
- [4] Subramanya R. Dulloor, Sanjay Kumar, Anil Keshavamurthy, Philip Lantz, Dheeraj Reddy, Rajesh Sankaran, and Jeff Jackson. 2014. System Software for Persistent Memory. In *Proceedings of the 9th European Conference on Computer Systems (EuroSys '14)*. ACM, Article 15, 15 pages.
- [5] IBM Corporation. 1978. *IBM System/38 Technical Developments*. Technical Report G580-0237-1. IBM.
- [6] Intel Corporation. 2023. *Intel 64 and IA-32 Architectures Software Developer's Manual, Volume 3A: System Programming Guide*. Intel Corporation.
- [7] Y. Nir and A. Langley. 2015. *ChaCha20 and Poly1305 for IETF Protocols*. RFC 7539. Internet Engineering Task Force.
- [8] NVM Express, Inc. 2021. *NVM Express Base Specification, Revision 2.0*. NVM Express, Inc.
- [9] Diego Ongaro and John Ousterhout. 2014. In Search of an Understandable Consensus Algorithm. In *2014 USENIX Annual Technical Conference (USENIX ATC '14)*. USENIX Association, 305–320.
- [10] Jian Xu and Steven Zhang. 2016. NOVA: A Log-structured File System for Hybrid Volatile/Non-volatile Main Memories. In *14th USENIX Conference on File and Storage Technologies (FAST '16)*. USENIX Association, 323–338.
- [11] Ziyi Yang, James R. Harris, Benjamin Walker, Daniel Verkamp, Changpeng Liu, Cunyin Chang, Gang Cao, Jonathan Stern, Vishal Verma, and Luse Paul. 2017. SPDK: A Development Kit to Build High Performance Storage Applications. In *2017 IEEE International Conference on Cloud Computing Technology and Science (CloudCom)*. IEEE, 154–161.